

DARK STORE DOWN

Zipto Quick Commerce — Noida Cluster Review

Hand this to students. It contains no answers.

The situation

Zipto runs ten 10-minute grocery dark stores across the Noida cluster. Most of them make money. Some do not.

You are the analytics team. The COO has given you six weeks of order, delivery and product data, and a board meeting on Monday.

"Which two stores do I fix first, what exactly is broken, and what is fixing it worth in rupees?"

You have until 4:00 PM.

Your team

Three people. **Three lanes. You each own one.**

Lane	You own these files	You answer
1 — Store Performance	dark_stores_messy.csv, store_hours_messy.csv	Which stores make money, and which lose it?
2 — Delivery & Fulfilment	trip_logs_messy.csv, delivery_partners_messy.csv	Are we delivering what we say we deliver?
3 — Product & Promotions	order_items_messy.csv, products_messy.csv, promo_codes_messy.csv	What are we selling, and what are the discounts costing us?

Everyone shares `orders_clean.csv` — the spine. It is already clean. **Your own files are not.** Cleaning them is part of your job and it is scored.

You can start at 9:45 without waiting for anyone.

The economics

Use exactly this definition. Do not invent your own — your three answers must add up.

delivered → contribution = net_amount – cogs_amount – delivery_cost

returned → contribution = –(cogs_amount + delivery_cost)

cancelled → contribution = 0

store P&L = SUM(contribution) – (monthly_rent × months in the data)

What you hand in at 4:00 PM sharp

1. **Three Finding Cards** — one per person, in your own name. Non-negotiable.
2. **One Cleaning Log** — everyone adds their own rows, as you go.
3. **One decision slide deck** — five slides maximum, or a dashboard.

Your decision must contain:

- **Top 2 stores** to fix
 - **Root cause** for each — the cause, not the symptom
 - **Evidence** — the numbers, and where they came from
 - **₹ impact per month** — *showing the arithmetic*
 - **Recommendation** — what the COO does on Monday
 - **What would make us wrong** — one honest line
-

How you will be judged

	Points
Analytical correctness	25
Business insight & recommendation	25
Communication & story	20
Cleaning rigour	15
Craft of the deliverable	10
Individual defence	5
<i>Technical stretch (Python / DAX / dashboards)</i>	<i>+5 bonus</i>

Read that again. Tooling is worth 5 bonus points out of 105. A sharp answer built in Excel and SQL will beat a mediocre one built in Power BI. There is no advantage to being further along in the course — there is an advantage to thinking clearly.

Four things that lose marks

- A ₹ figure you cannot rebuild live when a judge asks
 - A cleaning log written at 3:55 PM instead of as you went
 - Comparing stores on raw totals without checking whether that comparison is fair
 - Leaving "what would make us wrong" blank
-

The day

09:30	Brief and data drop
09:45	Lane work — clean your files, start hunting
11:30	Number Check — 10 minutes. Say your headline number out loud to your team. Bell rings.

11:40	Finish your Finding Card
13:00	Lunch
13:45	Buffer — finish up, or hunt for more
14:45	Consolidation — merge three cards into one decision
15:30	Build and polish
16:00	HARD FREEZE. No exceptions.
16:15	Judging pods — every one of you will be questioned on your own card
17:15	Top 3 present on the main stage
17:50	Awards

Stuck?

Each lane has a **sealed hint envelope**. Opening it costs your team 5 points and the facilitator will note it. It is there so you keep moving — use it if you are genuinely blocked for more than 30 minutes, but try everything else first.

Awards

- **Winning team**
- **Best Cleaning Log**
- **Best Individual Finding Card** — this one is yours alone, and it is the one to want.

Everything you produce today is yours to keep. Your Finding Card is a defensible piece of work you can put in front of an interviewer on Monday — so write it like someone will read it who was not in the room.